Computer Architecture
（计算机体系结构）

Lecturer
Yuanqing
Cheng

Lecture 23– CPU Design : Pipelining to Improve Performance

2020-10-23

Review: Single cycle datapath

- 5 steps to design a processor
  1. Analyze instruction set datapath requirements
  2. Select set of datapath components & establish clock methodology
  3. Assemble datapath meeting the requirements
  4. Analyze implementation of each instruction to determine setting of control points that effects the register transfer.
  5. Assemble the control logic
- Control is the hard part
- MIPS makes that easier
  - Instructions same size
  - Source registers always in same place
  - Immediates same size, location
  - Operations always on registers/ immediates

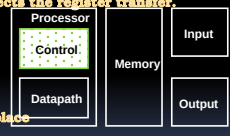

How We Build The Controller

RegDst = add + sub
ALUSrc = ori + lw + sw
MemtoReg = lw
RegWrite = add + sub + ori + lw
MemWrite = sw
nPCsel = beq
Jump = jump
ExtOp = lw + sw
ALUctr[0] = sub + beq (assume ALUctr is 0 ADD, 01: SUB, 10: OR)
ALUctr[1] = or

where,

rtype = $\sim op_5 \bullet \sim op_4 \bullet \sim op_3 \bullet \sim op_2 \bullet \sim op_1 \bullet \sim op_0$,
ori = $\sim op_5 \bullet \sim op_4 \bullet op_3 \bullet op_2 \bullet \sim op_1 \bullet op_0$
lw = $op_5 \bullet \sim op_4 \bullet \sim op_3 \bullet \sim op_2 \bullet op_1 \bullet op_0$
sw = $op_5 \bullet \sim op_4 \bullet op_3 \bullet \sim op_2 \bullet op_1 \bullet op_0$
beq = $\sim op_5 \bullet \sim op_4 \bullet \sim op_3 \bullet op_2 \bullet \sim op_1 \bullet \sim op_0$
jump = $\sim op_5 \bullet \sim op_4 \bullet \sim op_3 \bullet \sim op_2 \bullet op_1 \bullet \sim op_0$

add = rtype $\bullet$ $func_5 \bullet \sim func_4 \bullet \sim func_3 \bullet \sim func_2 \bullet \sim func_1 \bullet \sim func_0$
sub = rtype $\bullet$ $func_5 \bullet \sim func_4 \bullet \sim func_3 \bullet \sim func_2 \bullet func_1 \bullet \sim func_0$

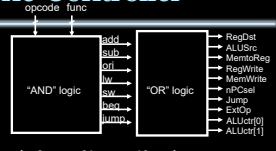

Omigosh omigosh, do you know what this means?

Call home, we've made HW/SW contact!

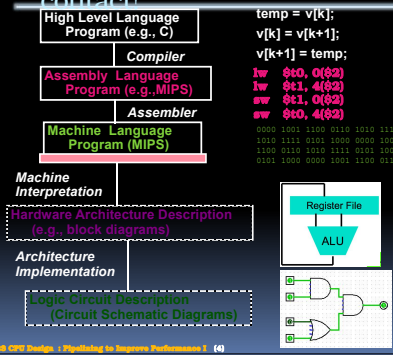


```
temp = v[k];
v[k] = v[k+1];
v[k+1] = temp;
lw  $t0, 0($t2)
lw  $t1, 4($t2)
sw  $t1, 0($t2)
sw  $t0, 4($t2)
0000 1001 1100 0110 1010 1111 0101 1000
1010 1111 0101 1000 0000 1001 1100 0110
1100 0110 1010 1111 0101 1000 0000 1001
0101 1000 0000 1001 1100 0110 1010 1111
```


Processor Performance

- Can we estimate the clock rate (frequency) of our single-cycle processor? We know:
  - 1 cycle per instruction
  - lw is the most demanding instruction.
  - Assume these delays for major pieces of the datapath:
    - Instr. Mem, ALU, Data Mem : 2ns each, regfile 1ns
    - Instruction execution requires: 2 + 1 + 2 + 2 + 1 = 8ns
    - ⇒ 125 MHz
- What can we do to improve clock rate?
- Will this improve performance as well?
  - We want increases in clock rate to result in programs executing quicker.

Gotta Do Laundry

- Ann, Brian, Cathy, Dave each have one load of clothes to wash, dry, fold, and put away
  - Washer takes 30 minutes
  - Dryer takes 30 minutes
  - "Folder" takes 30 minutes
  - "Stasher" takes 30 minutes to put clothes into drawers

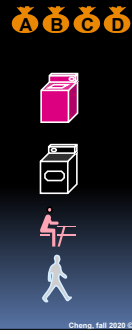

Sequential Laundry

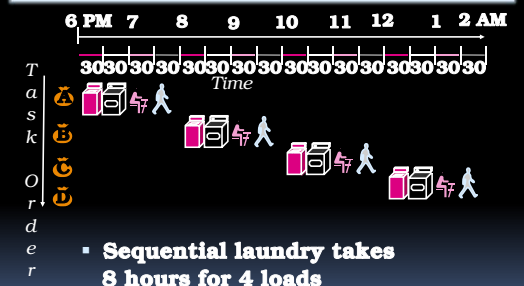

- Sequential laundry takes 8 hours for 4 loads

Pipelined Laundry

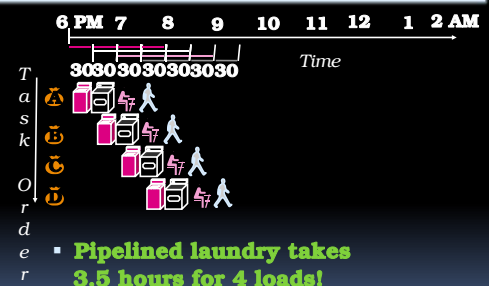

- Pipelined laundry takes 3.5 hours for 4 loads!

General Definitions

- **Latency**: time to completely execute a certain task
  - for example, time to read a sector from disk is disk access time or disk latency
- **Throughput**: amount of work that can be done over a period of time

Pipelining Lessons (1/2)

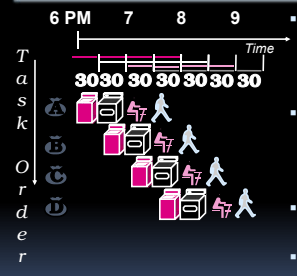

- Pipelining doesn't help **latency** of single task, it helps **throughput** of entire workload
- **Multiple** tasks operating simultaneously using different resources
- Potential speedup = **Number pipe stages**
- Time to "**fill**" pipeline and time to "**drain**" it reduces speedup: 2.3X v. 4X in this example

Pipelining Lessons (2/2)

- Suppose new Washer takes 20 minutes, new Stasher takes 20 minutes. How much faster is pipeline?
- Pipeline rate limited by **slowest** pipeline stage
- Unbalanced lengths of pipe

Steps in Executing MIPS

1) **IFtch**: Instruction Fetch, Increment PC
2) **Dcd**: Instruction Decode, Read Registers
3) **Exec**:
   Mem-ref: Calculate Address
   Arith-log: Perform Operation
4) **Mem**:
   Load: Read Data from Memory
   Store: Write Data to Memory
5) **WB**: Write Data Back to Register

Pipelined Execution Representation

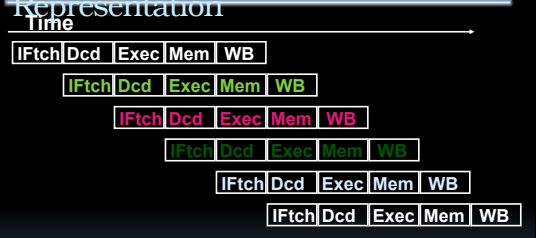

- **Every instruction must take same number of steps, also called pipeline "stages", so some will go idle sometimes**

Review: Datapath for MIPS

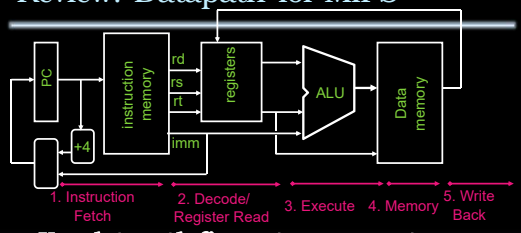

- **Use datapath figure to represent pipeline**

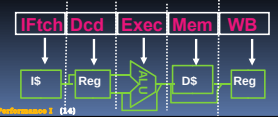

Graphical Pipeline Representation

(In Reg, right half highlight read, left half write)

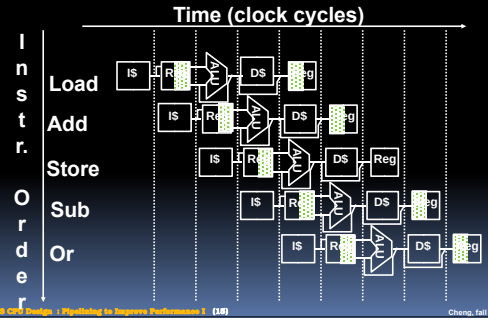

Example

- **Suppose 2 ns for memory access, 2 ns for ALU operation, and 1 ns for register file read or write; compute instruction rate**
- **Nonpipelined Execution:**
  - `lw` **: IF + Read Reg + ALU + Memory + Write Reg = 2 + 1 + 2 + 2 + 1 = 8 ns**
  - `add`**: IF + Read Reg + ALU + Write Reg = 2 + 1 + 2 + 1 = 6 ns**
    (recall 8ns for single-cycle processor)
- **Pipelined Execution:**
  - **Max(IF,Read Reg,ALU,Memory,Write Reg) = 2 ns**

Pipeline Hazard: Matching socks in later load

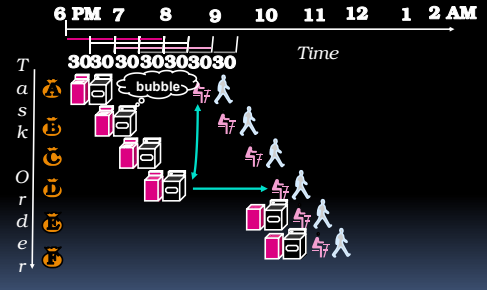

- **A depends on D; stall since folder tied up**

Problems for Pipelining CPUs

- **Limits to pipelining: Hazards prevent next instruction from executing during its designated clock cycle**
  - **Structural hazards : HW cannot support some combination of instructions (single person to fold and put clothes away)**
  - **Control hazards : Pipelining of branches causes later instruction fetches to wait for the result of the branch**
  - **Data hazards : Instruction depends on result of prior instruction still in the pipeline (missing sock)**
- **These might result in pipeline stalls or "bubbles" in the pipeline.**

Structural Hazard #1: Single Memory (1/2)

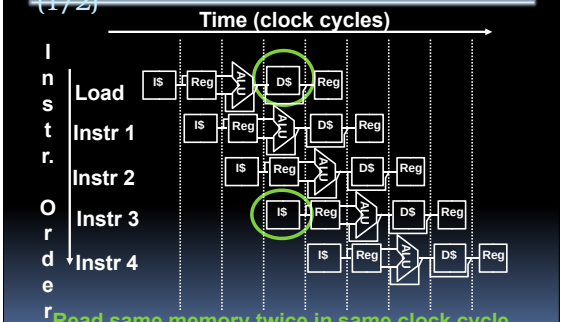

Read same memory twice in same clock cycle

Structural Hazard #1: Single Memory (2/2)

- **Solution:**
  - **infeasible and inefficient to create second memory**
  - **(We'll learn about this more next week)**
  - **so simulate this by having two Level 1 Caches (a temporary smaller [of usually most recently used] copy of memory)**
  - **have both an L1 Instruction Cache and an L1 Data Cache**
  - **need more complex hardware to control when both caches miss**

Structural Hazard #2: Registers (1/2)

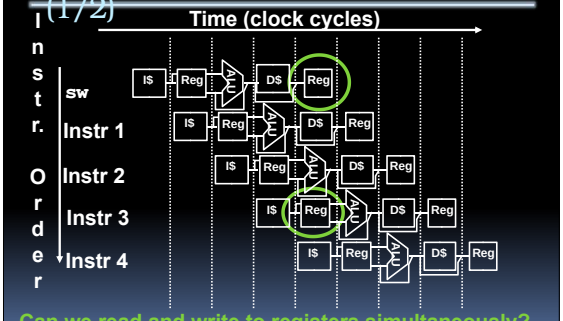

Can we read and write to registers simultaneously?

Structural Hazard #2: Registers (2/2)

- **Two different solutions have been used:**
  1) **RegFile access is *VERY* fast: takes less than half the time of ALU stage**
     - **Write to Registers during first half of each clock cycle**
     - **Read from Registers during second half of each clock cycle**
  2) **Build RegFile with independent read and write ports**
- **Result: can perform Read and Write during same clock cycle**

Peer Instruction

1) **Thanks to pipelining, I have reduced the time it took me to wash my one shirt.**
2) **Longer pipelines are always a win (since less work per stage & a faster clock).**

12
a) FF
b) FT
c) TF
d) TT

Peer Instruction Answer

1) **Throughput better, not execution time**
2) **"...longer pipelines do usually mean faster clock, but branches cause problems!"**

FALSE

FALSE

12
a) FF
b) FT
c) TF
d) TT

Things to Remember

- **Optimal Pipeline**
  - **Each stage is executing part of an instruction each clock cycle.**
  - **One instruction finishes during each clock cycle.**
  - **On average, execute far more quickly.**
- **What makes this work?**
  - **Similarities between instructions allow us to use same stages for all instructions (generally).**
  - **Each stage takes about the same amount of time as all others: little wasted time.**