
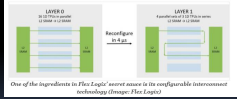


Computer Architecture  
(计算机体系结构)


**Lecture 25 – CPU Design :  
Pipelining to Improve  
Performance II**  
2020-10-26

Lecturer  
Yuangqing  
Cheng

Flex Logix' Edge AI Accelerator  
Battles Nvidia on Price-Performance

One of the ingredients in Flex Logix' secret sauce is its configurable interconnect technology (Stream Flex Interconnect).



Working together on-chip can be improved in the same time, without compromising the interconnect. Configurable interconnect can be increased (change the layout).

## Review

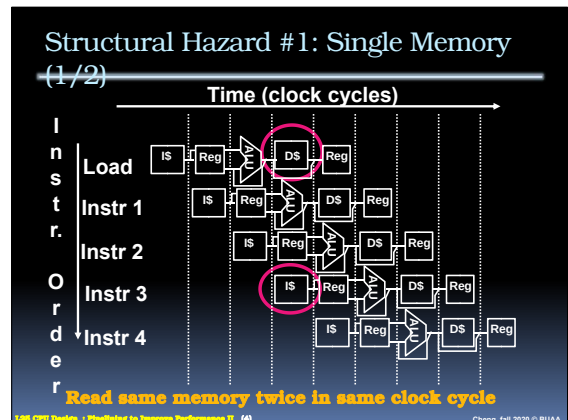
- **Pipelining is a BIG idea**
- **Optimal Pipeline**
  - **Each stage is executing part of an instruction each clock cycle.**
  - **One instruction finishes during each clock cycle.**
  - **On average, execute far more quickly.**
- **What makes this work?**
  - **Similarities between instructions allow us to use same stages for all instructions (generally).**
  - **Each stage takes about the same amount of time as all others. Little wasted time.**

L26 CPU Design | Pipelining to Improve Performance II (3)

## Problems for Pipelining CPUs

- **Limits to pipelining: Hazards prevent next instruction from executing during its designated clock cycle**
  - **Structural hazards** : HW cannot support some combination of instructions (single person to fold and put clothes away)
  - **Control hazards** : Pipelining of branches causes later instruction fetches to wait for the result of the branch
  - **Data hazards** : Instruction depends on result of prior instruction still in the pipeline (missing sock)
- **These might result in pipeline stalls or bubbles in the pipeline.**

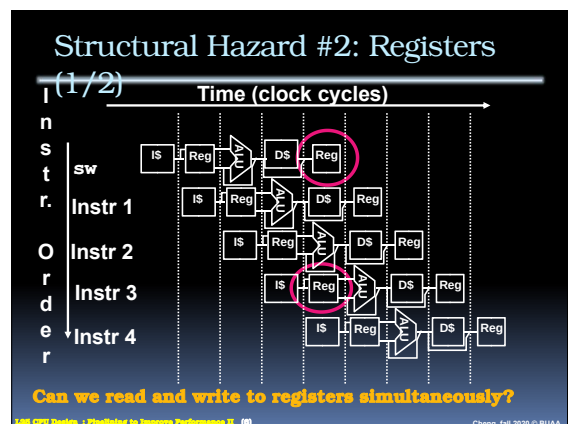
L26 CPU Design | Pipelining to Improve Performance II (4)



## Structural Hazard #1: Single Memory (2/2)

- **Solution:**
  - **infeasible and inefficient to create second memory**
  - **(We'll learn about this shortly)**
  - **...so simulate this by having two Level 1 Caches**
    - (a temporary smaller [of usually most recently used] copy of memory)
  - **have both an L1 Instruction Cache and an L1 Data Cache**
  - **need more complex hardware to control when both caches miss**

L26 CPU Design | Pipelining to Improve Performance II (5)



## Structural Hazard #2: Registers

(2/2)

- Two different solutions have been used:

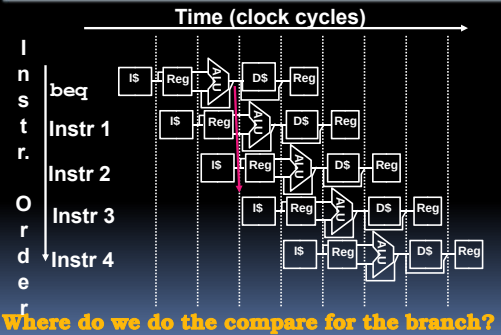
1) RegFile access is VERY fast: takes less than half the time of ALU stage

- Write to Registers during first half of each clock cycle
- Read from Registers during second half of each clock cycle

2) Build RegFile with independent read and write ports

- Result: can perform Read and Write during same clock cycle

## Control Hazard: Branching (1/9)



## Control Hazard: Branching (2/9)

- We had put branch decision-making hardware in ALU stage
  - therefore two more instructions after the branch will always be fetched, whether or not the branch is taken
- Desired functionality of a branch
  - if we do not take the branch, don't waste any time and continue executing normally
  - if we take the branch, don't execute any instructions after the branch, just go to the desired label

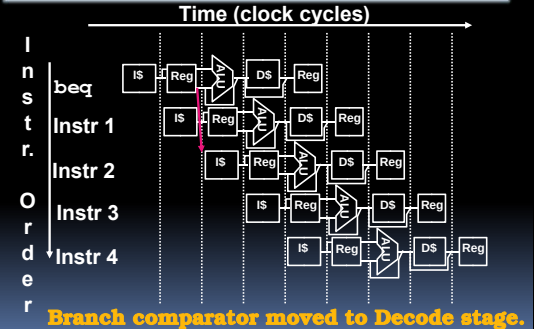
## Control Hazard: Branching (3/9)

- Initial Solution: Stall until decision is made
  - insert "no-op" instructions (those that accomplish nothing, just take time) or hold up the fetch of the next instruction (for 2 cycles).
  - Drawback: branches take 3 clock cycles each (assuming comparator is put in ALU stage)

## Control Hazard: Branching (4/9)

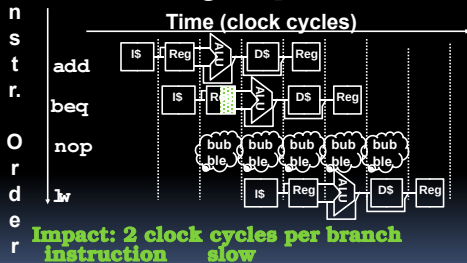
- Optimization #1:
  - insert special branch comparator in Stage 2
  - as soon as instruction is decoded ( Opcode identifies it as a branch), immediately make a decision and set the new value of the PC
  - Benefit: since branch is complete in Stage 2, only one unnecessary instruction is fetched, so only one no-op is needed
  - Side Note: This means that branches are idle in Stages 3, 4 and 5.

## Control Hazard: Branching (5/9)



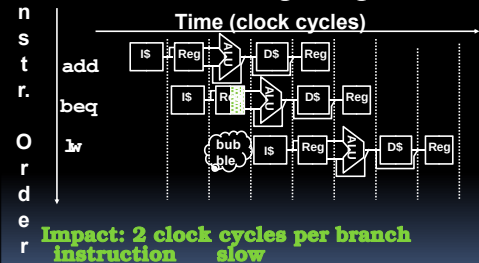
## Control Hazard: Branching (6/9)

### ▪ User inserting no-op instruction



## Control Hazard: Branching (7/9)

### ▪ Controller inserting a single bubble



...story about engineer, physicist, mathematician asked to build a fence around a flock of sheep using minimal fence...

## Control Hazard: Branching (8/9)

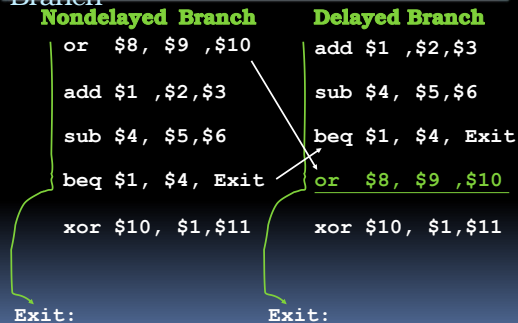
- **Optimization #2: Redefine branches**
  - Old definition: if we take the branch, none of the instructions after the branch get executed by accident
  - New definition: whether or not we take the branch, the single instruction immediately following the branch gets executed (called the **branch-delay slot**)
- The term **"Delayed Branch"** means we always execute inst after branch
- This optimization is used with MIPS

## Control Hazard: Branching (9/9)

- **Notes on Branch-Delay Slot**
  - **Worst-Case Scenario:** can always put a no-op in the branch-delay slot
  - **Better Case:** can find an instruction preceding the branch which can be placed in the branch-delay slot without affecting flow of the program
    - re-ordering instructions is a common method of speeding up programs
    - compiler must be very smart in order to find instructions to do this
    - usually can find such an instruction at least 50% of the time

Jumps also have a delay slot...

## Example: Nondelayed vs. Delayed Branch



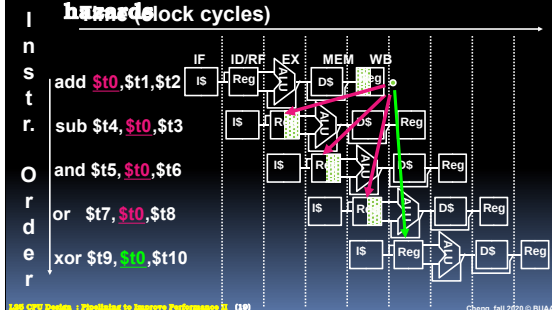
## Data Hazards (1/2)

### ▪ Consider the following sequence of instructions

```
add $t0, $t1, $t2
sub $t4, $t0, $t3
and $t5, $t0, $t6
or $t7, $t0, $t8
xor $t9, $t0, $t10
```


## Data Hazards (2/2)

- Data-flow backward in time are hazards (lock cycles)



## Data Hazard Solution:


- Forwarding
- Forward result from one stage to another



“or” hazard solved by register hardware

## Data Hazard: Loads (1/4)


- Dataflow backwards in time are hazards



- Can't solve all cases with forwarding
- Must stall instruction dependent on load, then forward (more hardware)

## Data Hazard: Loads (2/4)

- Hardware stalls pipeline
- Called “interlock”

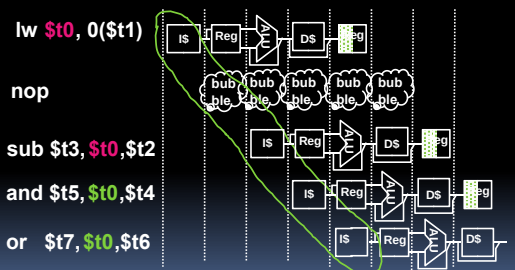


## Data Hazard: Loads (3/4)

- Instruction slot after a load is called “load delay slot”
- If that instruction uses the result of the load, then the hardware interlock will stall it for one cycle.
- If the compiler puts an unrelated instruction in that slot, then no stall
- Letting the hardware stall the instruction in the delay slot is equivalent to putting a nop in the slot (except the latter uses more code space)

## Data Hazard: Loads (4/4)

- Stall is equivalent to nop



## Peer Instruction

- 1) Thanks to pipelining, I have **reduced the time** it took me to wash my one shirt.
- 2) Longer pipelines are **always a win** (since less work per stage & a faster clock).

	12
a)	FF
b)	FT
c)	TF
d)	TT

## “And in Conclusion..”

- **Pipeline challenge is hazards**
  - Forwarding helps w/many data hazards
  - Delayed branch helps with control hazard in 5 stage pipeline
  - Load delay slot / interlock necessary
- **More aggressive performance:**
  - Superscalar
  - Out-of-order execution

## Bonus slides

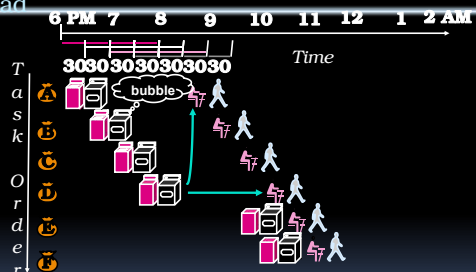
- These are extra slides that used to be included in lecture notes, but have been moved to this, the “bonus” area to serve as a supplement.
- The slides will appear in the order they would have in the normal presentation

# Bonus

## Historical Trivia


- First MIPS design did not interlock and stall on load-use data hazard
- Real reason for name behind MIPS: **Microprocessor without Interlocked Pipeline Stages**
  - Word Play on acronym for Millions of Instructions Per Second, also called MIPS

## Pipeline Hazard: Matching socks in later load



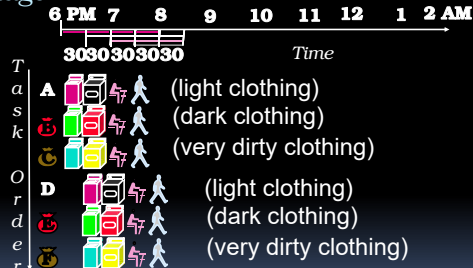
- A depends on D; stall since folder tied up; Note this is much different from processor cases so far. We have not had a earlier instruction depend on a later one.

## Out-of-Order Laundry: Don't Wait



- A depends on D; rest continue; need more resources to allow out-of-order

## Superscalar Laundry: Parallel per stage



- More resources, HW to match mix of parallel tasks?

## Superscalar Laundry: Mismatch



- Task mix underutilizes extra

## Peer Instruction (1/2)

Assume 1 instr/clock, delayed branch, 5 stage pipeline, forwarding, interlock on unresolved load hazards (after  $10^3$  loops, so pipeline full)

```
Loop: lw    $t0, 0($s1)
      addu  $t0, $t0, $s2
      sw    $t0, 0($s1)
      addiu $s1, $s1, -4
      bne   $s1, $zero, Loop
      nop
```

- How many pipeline stages (clock cycles) per loop iteration to execute this code?

1  
2  
3  
4  
5  
6  
7  
8  
9  
10

## Peer Instruction Answer (1/2)

- Assume 1 instr/clock, delayed branch, 5 stage pipeline, forwarding, interlock on unresolved load hazards.  $10^3$  iterations, so pipeline full hazard so stall)

```
Loop: 1. lw    $t0, 0($s1)
      2. addu  $t0, $t0, $s2
      3. sw    $t0, 0($s1)
      4. addiu $s1, $s1, -4
      5. bne   $s1, $zero, Loop
      6. nop   (delayed branch so exec. nop)
```

- How many pipeline stages (clock cycles) per loop iteration to execute this code?

1 2 3 4 5 6 7 8 9 10

## Peer Instruction (2/2)

Assume 1 instr/clock, delayed branch, 5 stage pipeline, forwarding, interlock on unresolved load hazards (after  $10^3$  loops, so pipeline full).

Rewrite this code to reduce pipeline stages (clock cycles) per loop to as few as possible.

```
Loop: lw    $t0, 0($s1)
      addu  $t0, $t0, $s2
      sw    $t0, 0($s1)
      addiu $s1, $s1, -4
      bne   $s1, $zero, Loop
      nop
```

- How many pipeline stages (clock cycles) per loop iteration to execute this code?

1  
2  
3  
4  
5  
6  
7  
8  
9  
10

## Peer Instruction (2/2) How long to execute this code to reduce clock cycles per loop to as few as possible:

```
Loop: 1. lw    $t0, 0($s1)
      2. addiu $s1, $s1, -4
      3. addu  $t0, $t0, $s2
      4. bne   $s1, $zero, Loop
      5. sw    $t0, +4($s1)
      (modified sw to put past addiu)
```

- How many pipeline stages (clock cycles) per loop iteration to execute your revised code? (assume pipeline is full)

1 2 3 4 5 6 7 8 9 10