



Computer Architecture (计算机体系结构)

Lecture 25 – Caches I 2020-10-26

Lecturer:
Yuangqing
Cheng


Review : Pipelining

- **Pipeline challenge is hazards**
 - **Forwarding helps w/many data hazards**
 - **Delayed branch helps with control hazard in our 5 stage pipeline**
 - **Data hazards w/Loads → Load Delay Slot**
 - **Interlock** → “smart” CPU has HW to detect if conflict with inst following load, if so it stalls
- **More aggressive performance (discussed in section next week)**
 - **Superscalar (parallelism)**
 - **Out-of-order execution**

L26 Outline I (3)

Cheng, fall 2020 © BUAA

The Big Picture



L26 Outline I (4)

Cheng, fall 2020 © BUAA

Memory Hierarchy

I.e., storage in computer systems

- **Processor**
 - **holds data in register file (~100 Bytes)**
 - **Registers accessed on nanosecond timescale**
- **Memory (we'll call "main memory")**
 - **More capacity than registers (~ Gbytes)**
 - **Access time ~50-100 ns**
 - **Hundreds of clock cycles per memory access?!**
- **Disk**
 - **HUGE capacity (virtually limitless)**
 - **VERY slow: runs ~milliseconds**


L26 Outline I (5)

Cheng, fall 2020 © BUAA

Motivation: Why We Use Caches

(written ☺)

- **1989 first Intel CPU with cache on chip**
- **1998 Pentium III has two cache levels on chip**



L26 Outline I (6)

Cheng, fall 2020 © BUAA

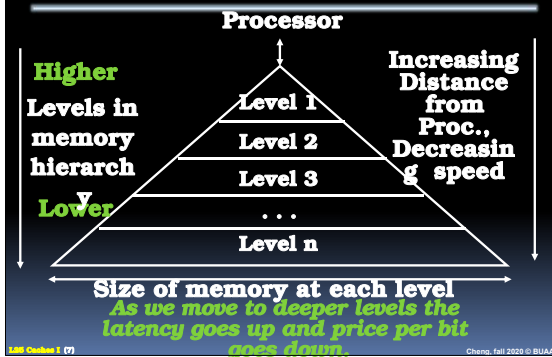
Memory Caching

- **Mismatch between processor and memory speeds leads us to add a new level: a memory **cache****
- **Implemented with same IC processing technology as the CPU (usually integrated on same chip): faster but more expensive than DRAM memory.**
- **Cache is a copy of a subset of main memory.**
- **Most processors have separate caches for instructions and data.**

L26 Outline I (7)

Cheng, fall 2020 © BUAA

Memory Hierarchy



Memory Hierarchy

- If level closer to Processor, it is:
 - **Smaller**
 - **Faster**
 - **More expensive**
 - **subset of lower levels (contains most recently used data)**
 - **Lowest Level (usually disk) contains all available data (does it go beyond the disk?)**
 - **Memory Hierarchy presents the processor with the illusion of a very large & fast memory**
- L200 Outline 1 (8) Cheng, fall 2020 © BUAA

Memory Hierarchy Analogy: Library (1/2)

- **You're writing a term paper (Processor) at a **table** in **Doe****
 - **Doe Library is equivalent to **disk****
 - **essentially limitless capacity**
 - **very slow to retrieve a book**
 - **Table is **main memory****
 - **smaller capacity: means you must return book when table fills up**
 - **easier and faster to find a book there once you've already retrieved it**
- L200 Outline 1 (9) Cheng, fall 2020 © BUAA

Memory Hierarchy Analogy: Library (2/2)

- **Open books on table are **cache****
 - **smaller capacity: can have very few open books fit on table; again, when table fills up, you must close a book**
 - **much, much faster to retrieve data**
 - **Illusion created: whole library open on the tabletop**
 - **Keep as many recently used books open on table as possible since likely to use again**
 - **Also keep as many books on table as possible, since faster than going to library**
- L200 Outline 1 (10) Cheng, fall 2020 © BUAA

Memory Hierarchy Basis

- **Cache contains copies of data in memory that are being used.**
 - **Memory contains copies of data on disk that are being used.**
 - **Caches work on the principles of **temporal and spatial locality**.**
 - **Temporal Locality: if we use it now, chances are we'll want to use it again soon.**
 - **Spatial Locality: if we use a piece of memory, chances are we'll use the neighboring pieces soon.**
- L200 Outline 1 (11) Cheng, fall 2020 © BUAA

Cache Design

- **How do we organize cache?**
 - **Where does each memory address map to?**
 - **(Remember that cache is subset of memory, so multiple memory addresses map to the same cache location.)**
 - **How do we know which elements are in cache?**
 - **How do we quickly locate them?**
- L200 Outline 1 (12) Cheng, fall 2020 © BUAA

Direct-Mapped Cache (1/4)

- In a **direct-mapped cache**, each **memory address is associated with one possible block** within the cache
 - Therefore, we only need to look in a **single location in the cache** for the data if it exists in the cache
 - Block is the unit of transfer** between cache and memory

L&S Chapter 1 (14)

Cheng, fall 2020 © BUAA

Direct-Mapped Cache

(2/4)

Memory Address Memory Index Cache 4 Byte Direct Mapped Cache



Block size = 1 byte

Cache Location 0 can

- be occupied by data from:
 - Memory location 0, 4, 8, ...
 - 4 blocks \Rightarrow any memory location must be bigger than one byte?

L&S Chapter 1 (14)

Cheng, fall 2020 © BUAA

Direct-Mapped Cache

(3/4)

Memory Address Memory Index Cache 8 Byte Direct Mapped Cache



Block size = 2 bytes

- When we ask for a byte, the system finds out the right block, and loads it all!
 - How does it know right block?
 - How do we select the byte?
- E.g., Mem address 11101?
 - How does it know WHICH colored block it originated from?
 - What do you do at baggage claim?

L&S Chapter 1 (15)

Cheng, fall 2020 © BUAA

Direct-Mapped Cache

(4/4)

Memory Address Memory Index Cache 8 Byte Direct Mapped Cache w/Tag!



Tag Data (Block size = 2 bytes)

- What should go in the tag?
 - Do we need the entire address?
 - What do all these tags have in common?
- What did we do with the immediate when we were branch addressing, always count by bytes?
- Why not count by **cache #**?
 - It's useful to draw memory with the same width as the block size

L&S Chapter 1 (16)

Cheng, fall 2020 © BUAA

Issues with Direct-Mapped

- Since **multiple memory addresses map to same cache index**, how do we tell which one is in there?
- What if we have a **block size > 1 byte**?
- Answer: divide memory address into **three fields**



L&S Chapter 1 (17)

Cheng, fall 2020 © BUAA

Direct-Mapped Cache

Terminology

- All fields are read as **unsigned integers**.
- Index**
 - specifies the cache index (which "row"/block of the cache we should look in)
- Offset**
 - once we've found correct block, specifies which byte within the block we want
- Tag**
 - the remaining bits after offset and index are determined; these are used to distinguish between all the memory addresses that map to the same location

L&S Chapter 1 (18)

Cheng, fall 2020 © BUAA


TIO great cache mnemonic

AREA (cache size, B)
= HEIGHT (# of blocks, $2^{(H+W)} = 2^H * 2^W$)

* WIDTH (size of one block, B/block)

Tag Index Offset

HEIGHT
(# of blocks)



L&R Question 1 (14)

Cheng, fall 2020 © BUAA

Direct-Mapped Cache Example

(1/3)

Suppose we have a 8B of data in a direct-mapped cache with 2 byte blocks

- Sound familiar?
- Determine the size of the tag, index and offset fields if we're using a 32-bit architecture
- Offset
 - need to specify correct byte within a block
 - block contains 2 bytes
= 2^1 bytes
 - need 1 bit to specify correct byte

L&R Question 1 (30)

Cheng, fall 2020 © BUAA

Direct-Mapped Cache Example

(2/3)

Index: (~index into an "array of blocks")

- need to specify correct block in cache
- cache contains 8 B = 2^3 bytes
- block contains 2 B = 2^1 bytes
- # blocks/cache
 - = $\frac{\text{bytes/cache}}{\text{bytes/block}}$
 - = $\frac{2^3 \text{ bytes/cache}}{2^1 \text{ bytes/block}}$
 - = 2^2 blocks/cache
- need 2 bits to specify this many blocks

L&R Question 1 (51)

Cheng, fall 2020 © BUAA

Direct-Mapped Cache Example

(3/3)

Tag: use remaining bits as tag

- tag length = addr length - offset - index
 - = $32 - 1 - 2$ bits
 - = 29 bits
- so tag is leftmost 29 bits of memory address
- Why not full 32 bit address as tag?
 - All bytes within block need same address (4b)
 - Index must be same for every address within a block, so it's redundant in tag check, thus can leave off to save memory (here 12 bits)

L&R Question 1 (59)

Cheng, fall 2020 © BUAA

And in Conclusion...

- We would like to have the capacity of disk at the speed of the processor: unfortunately this is not feasible.
- So we create a memory hierarchy:
 - each successively lower level contains "most used" data from next higher level
 - exploits temporal & spatial locality
 - do the common case fast, worry less about the exceptions (design principle of MIPS)
- Locality of reference is a Big Idea

L&R Question 1 (64)

Cheng, fall 2020 © BUAA